

Clickhouse_FETCH_PART

STOP MERGES

为MergeTree系列引擎表停止后台合并操作。

```
SYSTEM STOP MERGES [[db.]merge_tree_family_table_name]
```

stop insert 权限

1. 停止数据写入，确保迁移的数据的全量性
2. 停止merge 确保当前查到的所有数据 文件都没有变化

开始 fetch part

<https://github.com/ClickHouse/ClickHouse/pull/22706>

扩容需求

clickhouse 跨集群迁移

1 需求

cluster 1 数据 迁移值 cluster 2

2 解决方案

方案 1：clickhouse copier 生成配置文件

方案 2：clickhouse alter table fetch partition from "remote:/shardpath" 这里我们介绍方案2使用方法

3 源集群

1 表信息

```
BJHTYD-Hope-17-5.hadoop.jd.local :) select * from system.macros;
```

```
SELECT *
FROM system.macros
```

macro	substitution
replica	00
replica_dict	010-00
shard	010
shard_dict	1

```
show create table xxxx;
```

```
SHOW CREATE TABLE xxxx
```

```

└─statement─┘
|
| CREATE TABLE pc.xxxx
| (
|     `sale_mode` String,
|     `shop_id` String,
|     `dept_list` String,
|     `dt` Date,
|     `hash_key_for_shard` Int32
| )
ENGINE = ReplicatedMergeTree('/clickhouse/tables/pc/xxxx/{shard}', '{replica}')
PARTITION BY dt
ORDER BY (sale_mode, shop_id, dept_list)
SETTINGS storage_policy = 'jdob_ha', index_granularity = 8192

```

所以这里：当前shard 在zookeeper 中创建路径为： /clickhouse/tables/pc/xxxx/010

2 分区数据

```
select distinct partition from system.parts where table='xxxx';
```

```

SELECT DISTINCT partition
FROM system.parts
WHERE table = 'xxxx'

```

```

└─partition─┘
| 2020-09-16 |
| 2020-09-17 |
| 2020-09-18 |
| 2020-09-19 |
| 2020-09-20 |
| 2020-09-21 |
| 2020-09-22 |
| 2020-09-23 |
| 2020-09-24 |
| 2020-09-25 |
| 2020-09-26 |
| 2020-09-27 |
| 2020-09-28 |
| 2020-09-29 |
| 2020-09-30 |

```

| 2020-10-01 |

我们这里选择一个分区做测试

3 分区数据量

```
select count(*) from xxxx where dt='2020-09-16';
```

```
SELECT count(*)  
FROM xxx  
WHERE dt = '2020-09-16'
```

count()
16244

当前分区 为 '2020-09-16, 16244'

```
select count(*) from xxx where dt='2020-09-16';
```

```
SELECT count(*)  
FROM xxx  
WHERE dt = '2020-09-16'
```

count()
639228

```
1 rows in set. Elapsed: 2.347 sec. Processed 639.23 thousand rows, 1.28 MB  
(272.32 thousand rows/s., 544.64 KB/s.)
```

4 目标集群

·1 config.xml 源集群 zookeeper 集群添加

```
<auxiliary_zookeepers>
  <cluster1>
    <node index="1">
      <host></host>
      <port>2281</port>
    </node>
    <node index="2">
      <host></host>
      <port>2281</port>
    </node>
    <node index="3">
      <host></host>
      <port>2281</port>
    </node>
  </cluster1>
</auxiliary_zookeepers>
```

2 重启服务

System restart clickhouse-server*

3 fetch

```
ALTER TABLE xxx FETCH PARTITION '2020-09-16' FROM
'xxx:/clickhouse/tables/pc/xxx/010';
```

为什么能够支持 请查看我上一篇文章，这里主要是副本间通信，默认没有身份验证

```

//节点之间通信是不需要用户名和密码的
if (config().has("interserver_http_credentials"))
{
    String user = config().getString("interserver_http_credentials.user",
    "");
    String password =
    config().getString("interserver_http_credentials.password", "");

    if (user.empty())
        throw Exception("Configuration parameter interserver_http_credentials
user can't be empty", ErrorCodes::NO_ELEMENTS_IN_CONFIG);

    global_context->setInterserverCredentials(user, password);
}

```

interserver_http_credentials: 在使用Replicated *引擎进行复制期间进行身份验证的用户名和密码。这些凭据仅用于副本之间的通信，与ClickHouse客户端的凭据无关。服务器正在检查这些凭据以连接副本，并在连接到其他副本时使用相同的凭据。因此，对于群集中的所有副本，应将这些凭据设置为相同。默认不使用身份验证。

```

<interserver_http_credentials>
  <user>admin</user>
  <password>222</password>
</interserver_http_credentials>

```

操作日志

```

2021.04.02 15:10:14.286751 [ 145145 ] {3790665d-1aff-47d8-bef7-a5a718a9f97c}
<Debug> executeQuery: (from [::ffff:172.18.160.19]:52662) ALTER TABLE xxxx FETCH
PARTITION '2020-09-16' FROM 'HT0_CK_Pub_01:/clickhouse/tables/pc/xxxx/010';
2021.04.02 15:10:14.286860 [ 145145 ] {3790665d-1aff-47d8-bef7-a5a718a9f97c}
<Trace> ContextAccess (default): Access granted: ALTER FETCH PARTITION ON
jason.xxxx
2021.04.02 15:10:14.289740 [ 145145 ] {3790665d-1aff-47d8-bef7-a5a718a9f97c}
<Information> jason.xxxx (5352004e-b633-40f1-9b6b-fc064b5034bd): Will fetch
partition 20200916 from shard HT0_CK_Pub_01:/clickhouse/tables/pc/xxxx/010
(auxiliary zookeeper 'HT0_CK_Pub_01')
2021.04.02 15:10:14.297859 [ 145145 ] {3790665d-1aff-47d8-bef7-a5a718a9f97c}
<Information> jason.xxxx (5352004e-b633-40f1-9b6b-fc064b5034bd): Found 3
replicas, 3 of them are active. Selected 00 to fetch from.
2021.04.02 15:10:14.299013 [ 145145 ] {3790665d-1aff-47d8-bef7-a5a718a9f97c}
<Information> jason.xxxx (5352004e-b633-40f1-9b6b-fc064b5034bd): Parts to fetch:
1
2021.04.02 15:10:14.299034 [ 145145 ] {3790665d-1aff-47d8-bef7-a5a718a9f97c}
<Debug> jason.xxxx (5352004e-b633-40f1-9b6b-fc064b5034bd): Fetching part
20200916_1_1_0 from /clickhouse/tables/pc/xxxx/010/replicas/00
2021.04.02 15:10:14.300032 [ 145145 ] {3790665d-1aff-47d8-bef7-a5a718a9f97c}
<Trace> ReadWriteBufferFromHTTP: Sending request to http://10.198.17.5:9509/?
endpoint=DataPartsExchange%3A%2Fclickhouse%2Ftables%2Fpc%2Fxxxx%2F010%2Freplicas%
2F00&part=20200916_1_1_0&client_protocol_version=5&compress=false
2021.04.02 15:10:14.344212 [ 145145 ] {3790665d-1aff-47d8-bef7-a5a718a9f97c}
<Debug> jason.xxxx (5352004e-b633-40f1-9b6b-fc064b5034bd): Fetched part
20200916_1_1_0 from /clickhouse/tables/pc/xxxx/010/replicas/00 (to 'detached'
directory)
2021.04.02 15:10:21.870276 [ 145119 ] {} <Debug> jason.xxxx
(ReplicatedMergeTreeRestartingThread): Activating replica.
2021.04.02 15:10:21.872309 [ 145119 ] {} <Debug> jason.xxxx
(ReplicatedMergeTreeQueue): Loading queue from
/clickhouse/tables/jason/xxxx/00/replicas/01/queue
2021.04.02 15:10:21.872815 [ 145119 ] {} <Debug> jason.xxxx
(ReplicatedMergeTreeQueue): Having 0 queue entries to load, 0 entries already
loaded.
2021.04.02 15:10:21.873127 [ 145119 ] {} <Trace> jason.xxxx
(ReplicatedMergeTreeQueue): Loaded queue

```

fetch 后数据不可见 在 detach 目录存放

4 attach partition

```
alter table xxxx attach partition '2020-09-16';
```

5 数据校验

```
select count(*) from xxxx;
```

```
SELECT count(*)  
FROM xxxx
```

count()
16244

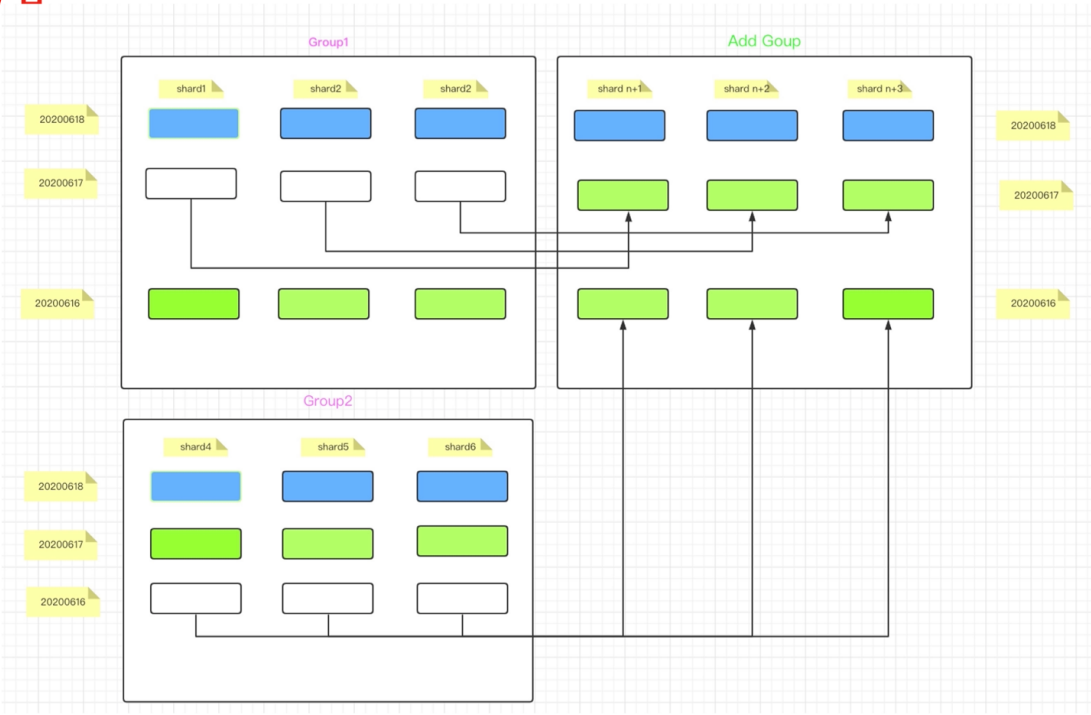
```
1 rows in set. Elapsed: 0.005 sec.
```

5 完成

6 后续优化

1. auxiliary_zookeepers 动态更新
2. 支持更细力度的 如 fetach part
3. 目前应该是 不支持相同part 名字落入同一个目前集群，如果支持
4. 直接支持集群迁移的命令

迁移-扩容



京东